

Folk Models of Good and Bad Taste: A Computational Text Analysis

Omar Lizardo
Department of Sociology, UCLA

July 20, 2026

Abstract

Sociologists have long debated the nature of “good” and “bad” taste, yet how lay individuals implicitly define these categories remains understudied. Using open-ended survey responses, this research note applies Non-Negative Matrix Factorization (NMF) to inductively map folk cultural models of taste definitions. I identify four distinct cultural models for “good taste” (Shared Taste, Cultivated Discernment, Personal Style, Aesthetic Curation) and “bad taste” (Relational Distance, Questionable Consumption, Lack of Refinement, Aesthetic Neglect). Chi-square cross-tabulations demonstrate that these models align into coherent, overarching cultural logics linking positive and negative evaluations. Furthermore, multinomial analyses show that these models are systematically patterned by socio-demographic factors, particularly gender and age. Finally, domain-mapping heatmaps across 19 cultural areas reveal that while core domains like Music and Clothing are universally subject to taste distinctions, individuals holding aesthetic and cultivated models exhibit a significantly higher overall propensity to draw taste boundaries across non-core domains.

1 Introduction

The sociological study of cultural stratification has traditionally focused on how cultural consumption patterns map onto social hierarchies (Bourdieu, 1984). While recent interview-based work has begun to explore everyday definitions of taste (Haak and Wilterdink, 2019; Jarness and Friedman, 2017), sociologists rarely measure how lay individuals systematically conceptualize “good” and “bad” taste. Do people view good taste as an acquired aesthetic skill, personal self-expression, or shared consensus with peers? Conversely, do they define bad taste as social or distance from the self, a lack of cultivation, or a passive neglect of aesthetic standards? Uncovering these implicit folk models provides crucial insight into contemporary mechanisms of symbolic boundary-drawing and cultural stratification.

2 Research Questions

This research note is guided by four primary research questions:

1. What are the distinct, latent cultural models that lay individuals use to define “good taste” and “bad taste”?
2. How do cultural models of good taste cross-tabulate with cultural models of bad taste to form internally coherent, broader cultural logics?
3. Are these cultural models of taste systematically patterned by traditional socio-demographic hierarchies (e.g., age, gender, and education)?
4. Does the propensity to draw boundaries (the total volume of differentiated cultural domains) vary significantly depending on the underlying cultural model an individual adopts?

3 Methods

3.1 Data

The data for this analysis are drawn from a study by Bonard et al. (2024) investigating folk conceptions of aesthetic taste. The original authors recruited a representative sample of the US population through the online participant platform Prolific Academic, compensating respondents £2.38 for their participation. A total of 297 individuals initially completed the survey. In accordance with data quality protocols, 56 respondents were excluded for failing to pass four embedded attention checks. This yielded a final analytic sample of 241 participants ($M_{age} = 44.92, SD = 16.12$; 124 men, 114 women, and 3 identifying as ‘other’). For demographic modeling, missing ages and unrealistic outliers (under 18) were safely dropped listwise. Additionally, to prevent model separation due to small sample sizes, extreme upper bounds of education (PhDs) were collapsed into a general graduate degree category, while parent’s education was dropped from the final models due to structural missingness (85% missing data owing to a survey skip logic error).

The survey instrument asked respondents whether they generally distinguish between good and bad taste. Those who answered affirmatively were prompted to provide open-ended textual definitions of “good taste” and “bad taste.” These open-access, publicly available responses serve as the primary corpus for the present computational analysis.

3.2 Text Analysis

Among the 241 participants, 224 (93%) indicated that they make a distinction between good and bad taste and thus provided open-ended definitions. These definitions were generally brief but highly informative. For “good taste,” responses averaged 170.15 characters in length ($SD = 90.41$, range: 26–469). For “bad taste,” definitions were slightly shorter, averaging 155.83 characters ($SD = 89.78$, range: 18–647). Combined, the average participant contributed 325.98 total characters of textual data ($SD = 166.99$, range: 45–899).

To analyze these short, unstructured participant-level text responses (hereafter “documents”), I used **Non-Negative Matrix Factorization (NMF)**. NMF is a linear-algebraic technique that reduces high-dimensional text data into a lower-dimensional set of latent “topics” that are seen as generative of the word choices in each document (Pauca et al., 2004; Xu et al., 2003). This approach has been found to be particularly effective for clustering short documents (Yan et al., 2013; Shi et al., 2018), such as the ones considered here, which consist of just a few sentences per respondent. First, the raw text was converted into an $N \times M$ Document-Term Matrix (V), where N is the number of documents and M is the size of the vocabulary. The values in this matrix were computed using Term Frequency-Inverse Document Frequency (TF-IDF), which weights words by how frequently they appear in a specific document while discounting words that are overwhelmingly common across the entire corpus. Specifically, the TF-IDF weight for a term t in a document d within a corpus D is calculated as:

$$\text{TF-IDF}(t, d, D) = \text{tf}(t, d) \cdot \log \left(\frac{N}{|\{d \in D : t \in d\}|} \right) \quad (1)$$

where $\text{tf}(t, d)$ is the raw count of term t in document d , and the logarithmic term represents the inverse document frequency. Within this logarithmic term, N is the total number of documents, and the denominator $|\{d \in D : t \in d\}|$ is the absolute count of documents within the entire corpus that contain the term t at least once. This ensures that terms appearing in nearly every document receive a weight close to zero, highlighting only the specific lexical signals that distinguish sub-populations. During preprocessing, domain-specific survey artifacts (e.g., “good,” “bad,” “taste,” “don’t”) and standard English stop words were removed.

The NMF algorithm then approximates V by factoring it into two lower-rank, strictly non-negative matrices: W and H (such that $V \approx WH$). The W matrix (size $N \times K$) represents the document-topic associations, providing a continuous “soft-assignment” weight of how strongly each respondent’s answer aligns with each of the K latent cultural models. The H matrix (size $K \times M$) represents the topic-term associations, indicating the relative importance of each vocabulary word in defining a given cultural model. The algorithm was optimized for each of the two text fields independently to prevent vocabulary blurring.

Finally, to explore the structural relationships between the documents and the derived topics, I conducted a Correspondence Analysis (CA). The CA was computed directly on the $N \times K$ soft-assignment matrix of document-to-topic probabilities generated by the NMF (with a minimal constant added to prevent division-by-zero errors), visualizing the geometric distance between respondents’ nuanced linguistic choices across the components.

3.3 Parameter Selection and Model Robustness

To ensure our topic solutions were empirically robust, we evaluated NMF models across a range of component configurations (K). The optimal configuration balances two competing demands: maximizing the granularity of semantic clusters while minimizing reconstruction error. We optimized for a 4-topic solution for both Good Taste and Bad Taste. This configuration prevents the model from collapsing into overly generic models while maintaining distinct qualitative variation.

4 Results

4.1 Cultural Models of Taste

Figure 2 and Figure 3 display the top ten most distinguishing terms (weighted by their NMF prominence) associated with each model, while Figure 1 shows the sample prevalence of each model. As shown in the term charts, the models are anchored by distinct lexical signatures that reveal different conceptualizations of aesthetic boundary-drawing.

For **Good Taste**, four primary cultural models emerged from the NMF analysis: *Shared Taste* (highlighting subjective similarity, driven by terms like “similar,” “say,” and “usually”), *Cultivated Discernment* (focusing on elevated aesthetic appreciation and driven by words like “art,” “quality,” and “appreciate”), *Personal Style* (emphasizing visual presentation and curation, characterized by words like “look,” “appealing,” and “clothing”), and *Aesthetic Curation* (centering on innate aesthetic stylishness via terms like “pleasing,” “aesthetically,” and “choices”). Table 1 and Table 2 present the term arrays and representative qualitative quotes for these models.

Topic	Count	Name	Representation
0	46	Shared Taste	similar, say, usually, likes, opinion, interests, certain...
1	99	Cultivated Discernment	art, appreciate, quality, enjoy, food, way, style...
2	39	Personal Style	look, appealing, clothing, way, knows, pick, listen...
3	40	Aesthetic Curation	pleasing, choices, aesthetically, make, stylish, typically, ...

Table 1: Good Taste Def Summary

For **Bad Taste**, respondents defined it through four distinct cultural models that mirror and invert the positive models: *Relational Distance* (emphasizing subjective differences and disagreement via words like “likes,” “different,” and “opinion”), *Questionable Consumption* (focusing on poor specific choices in domains like music and clothing), *Lack of Refinement*

Topic	Name	Representative Document
0	Shared Taste	“I’d say a person has good taste if I like similar things.”
1	Cultivated Discernment	“What I mean by that is they have a taste similar to mine or one I can appreciate. People who enjoy the same type of books or movies or art. Though I can understand someone may have exceptional taste and I don’t agree with it. Such as enjoying Rothko’s art or Wharton’s novels. ”
2	Personal Style	“Something that they like has an appealing look”
3	Aesthetic Curation	“Their choices are aesthetically pleasing to me. ”

Table 2: Representative Quotes for Good Taste Cultural Models

(centering on low quality, cheapness, or ugliness), and *Aesthetic Neglect* (highlighting gaudiness, clashing styles, and a general lack of care in appearance or curation, anchored by words like “care,” “really,” and “look”). Table 3 and Table 4 present the corresponding term summaries and representative quotes.

Topic	Count	Name	Representation
0	62	Relational Distance	likes, different, opinion, say, art, dislike, interests...
1	66	Questionable Consumption	choices, music, clothing, preferences, shows, food, unappeal...
2	50	Lack of Refinement	usually, quality, ugly, poor, way, likes, looks...
3	46	Aesthetic Neglect	care, really, look, colors, items, gaudy, know...

Table 3: Bad Taste Def Summary

Topic	Name	Representative Document
0	Relational Distance	“When I say that someone has bad taste, it means that in my opinion, the things that person likes is very different from the things I like, so, therefore, they have bad taste.”
1	Questionable Consumption	“Someone who has bad taste in clothing choices, music, food, dating preferences. ”
2	Lack of Refinement	“Usually it is about gaining notoriety at the expense of quality. It is a the immature desire tp shock others in negative ways and usually means poor quality.”
3	Aesthetic Neglect	“Someone who has bad taste doesn’t really care or even notice how an object is made. It could be junky and poorly finished. This is like an old dented car that was painted by a kid with a spray can of some garish color paint.”

Table 4: Representative Quotes for Bad Taste Cultural Models

4.2 Cross-Cultural Model Logics

To test whether these cultural models are isolated or form coherent “logics,” I cross-tabulated respondents’ primary cultural models across the two questions. The associations are statistically significant ($\chi^2 = 20.88, p < 0.05$). For instance, individuals who define good taste as




Figure 1: Topic Frequencies for Good and Bad Taste (Definitions)




Figure 2: Top Keywords for Good Taste Cultural Models




Figure 3: Top Keywords for Bad Taste Cultural Models

subjective similarity (*Shared Taste*) strongly tend to define bad taste as subjective difference (*Relational Distance*). Likewise, those focused on aesthetics (*Aesthetic Curation*) often define bad taste through the lens of specific choices (*Questionable Consumption*).

4.3 Demographics

To examine whether taste cultural models are patterned by socio-demographic factors, I fit multinomial logistic regression models predicting respondents' primary topic cultural model using Age, Gender, Education Level, and Political Orientation. The Likelihood Ratio χ^2 tests are presented in Table 5.

As shown in Table 5, sociodemographic background—particularly age and gender—structures how respondents conceptualize taste. For the Good Taste models, Age ($p \approx 0.05$) and Gender ($p < 0.05$) are significant predictors. The pattern is even stronger for Bad Taste models, where both Age ($p < 0.01$) and Gender ($p < 0.05$) significantly predict topic assignment. Education and Political Ideology did not reach statistical significance in these robust models, suggesting that generational cohorts and gender play a more primary role than formal educational attainment in shaping folk definitions of taste.

To unpack these significant effects, I generated predicted probabilities with 95% confidence intervals for both age and gender holding other variables constant. Figure 7 and Figure 8 display the faceted age trajectories for Good and Bad Taste models. A clear generational divide emerges: younger respondents are significantly more likely to conceptualize bad taste subjectively and relationally (e.g., *Likes Different Opinion*), whereas older adults shift heav-




Figure 4: Correspondence Analysis of Good Taste Cultural Models (Definitions)

Model	Predictor	LR χ^2	Df	<i>p</i> -value
Good Taste Def	Age (Polynomial)	12.44	6	0.0528
Good Taste Def	Gender	8.62	3	0.0348
Good Taste Def	Education Level	8.62	9	0.4729
Good Taste Def	Political Ideology	1.50	3	0.6832
Bad Taste Def	Age (Polynomial)	21.93	6	0.0012
Bad Taste Def	Gender	9.47	3	0.0236
Bad Taste Def	Education Level	11.93	9	0.2175
Bad Taste Def	Political Ideology	3.73	3	0.2917

Table 5: Multinomial Logistic Regression Likelihood Ratio Tests for Topic Predictors




Figure 5: Correspondence Analysis of Bad Taste Cultural Models (Definitions)




Figure 6: Pearson Residuals: Cross-Tabulation of Good and Bad Taste Cultural Models

ily toward objective, aesthetic judgments, overwhelmingly defining bad taste as a *Lack of Refinement* (e.g., “usually quality ugly”) and good taste through *Aesthetic Curation*.

Figure 9 and Figure 10 illustrate the predicted probabilities by gender. For good taste, women are significantly more likely than men to define it as a shared relational property (*Shared Taste*), while men are more likely to define it via pure aesthetics or highbrow discernment (*Art Appreciate Quality*). For bad taste, men are heavily overrepresented in the *Questionable Consumption* model, indicating a higher propensity to define bad taste simply by the consumption of the “wrong” things in specific domains like music and clothing.




Figure 7: Predicted Probabilities of Good Taste Models by Age (with 95% CIs)

4.4 Boundary Drawing and Domain Distinctions

To assess how broadly respondents apply their taste definitions across specific cultural domains, respondents were asked to rate whether it “makes sense” to distinguish good and bad taste across 19 different domains (from -3 = Strongly Disagree to +3 = Strongly Agree).

First, I constructed an additive variable representing the total volume of domains (out of 19) in which a respondent reported that making taste distinctions made sense (a score of 2 or above). I ran a one-way ANOVA to test whether the average number of differentiated domains varies depending on the respondent’s topic cultural model assignment. The results

Predicted Probabilities: Bad Taste Cultural Models
With 95% Confidence Intervals




Figure 8: Predicted Probabilities of Bad Taste Models by Age (with 95% CIs)




Figure 9: Predicted Probabilities of Good Taste Models by Gender




Figure 10: Predicted Probabilities of Bad Taste Models by Gender

demonstrate that respondents operating under different cultural models differ significantly in the sheer volume of cultural spaces they subject to aesthetic differentiation. Specifically, this relationship is robust for the *bad taste* cultural models (see Figure 11).




Figure 11: Mean Number of Distinguishable Domains by Bad Taste Cultural Model

Predictor	<i>F</i>	num Df	den Df	<i>p</i> -value
Good Taste Def	0.67	3	203	0.571
Bad Taste Def	5.64	3	203	0.001***

* $p < 0.05$; ** $p < 0.01$; *** $p < 0.001$

Table 6: ANOVA Results: Number of Domains (score ≥ 2) by Taste Cultural Model.

Second, visualizing the specific probabilities of boundary-drawing across these 19 domains via heatmaps (Figure 12) reveals several interesting patterns. Regardless of the cultural model adopted, respondents display a remarkably high probability of drawing taste distinctions in core lifestyle and aesthetic domains like **Clothing, Interior Design, Movies, Music, and Paintings**. Conversely, taste distinctions are almost universally rejected for domains like **Car Races and Wrestling**.

5 Discussion

This text-analytic exploration demonstrates that lay definitions of taste are not monolithic; they are multi-dimensional, spanning aesthetic, subjective, and quality-based dimensions. Crucially, these dimensions are symmetrical. If an individual views good taste through the lens of objective aesthetic criteria, they evaluate bad taste using the exact same yardstick.




Figure 12: Probability of Making Taste Distinctions by Domain and Bad Taste Cultural Model

Furthermore, by mapping these inductively derived cultural models back onto structured survey variables, we observe that the specific folk model an individual uses is significantly structured by age and gender, and in turn dictates the breadth of their cultural boundary-drawing. Future research should expand on these findings by examining how these deeply held, qualitative cultural models of taste interact with actual patterns of structural inequality.

References

- Bonard, C., Cova, F., and Humbert-Droz, S. (2024). De gustibus est disputandum: an empirical investigation of the folk concept of good taste. *Dataset*.
- Bourdieu, P. (1984). *Distinction: A Social Critique of the Judgement of Taste*. Harvard University Press, Cambridge, MA.
- Haak, M. v. d. and Wilterdink, N. (2019). Struggling with distinction: How and why people switch between cultural hierarchy and equality. *European Journal of Cultural Studies*, 22(4):416–432.
- Jarness, V. and Friedman, S. (2017). ‘i’m not a snob, but...’: Class boundaries and the downplaying of difference. *Poetics*, 61:14–25.
- Pauca, V. P., Shahnaz, F., Berry, M. W., and Plemmons, R. J. (2004). Text mining using non-negative matrix factorizations. In *Proceedings of the 2004 SIAM International Conference on Data Mining*, pages 452–456.
- Shi, T., Kang, K., Choo, J., and Reddy, C. K. (2018). Short-text topic modeling via non-negative matrix factorization enriched with local word-context correlations. In *Proceedings of the 2018 World Wide Web Conference*, pages 1105–1114.
- Xu, W., Liu, X., and Gong, Y. (2003). Document clustering based on non-negative matrix factorization. In *Proceedings of the 26th Annual International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 267–273.
- Yan, X., Guo, J., Liu, S., and Cheng, X. (2013). Learning topics in short texts by non-negative matrix factorization on term correlation matrix. In *Proceedings of the 2013 SIAM International Conference on Data Mining*, pages 749–757.